



## **NPTEL ONLINE CERTIFICATION COURSES**

**Operating System Fundamentals**

**Santanu Chattopadhyay**

**Electronics and Electrical Communication Engg.**

**Memory Management**

## CONCEPTS COVERED

### Concepts Covered:

- Background
- Swapping
- Contiguous Memory Allocation
- Segmentation
- Paging
- Structure of the Page Table
- Example Architectures



# Objectives

- To provide a detailed description of various ways of organizing memory hardware
- To discuss various memory-management techniques, including paging and segmentation
- To provide a detailed description of memory management policies in different architectures



# Background

- A program must be brought (from disk) into memory and placed within a process for it to be run
- A program can be written in machine language, assembly language, or high-level language.
- Main memory and registers are the only storage entities that a CPU can access directly
- The CPU fetches instructions from main memory according to the value of the program counter.
- Typical instruction execution cycle – fetch instruction from memory, decode the instruction, operand fetch, possible storage of result in memory.



## Background (Cont.)

- Memory unit only sees a stream of one of the following:
  - address + read requests (e.g., load memory location 20010 into register number 8)
  - address + data and write requests (e.g., store content of register 6 into memory location 1090)
- Memory unit does not know how these addresses were generated
- Register access can be done in one CPU clock (or less)



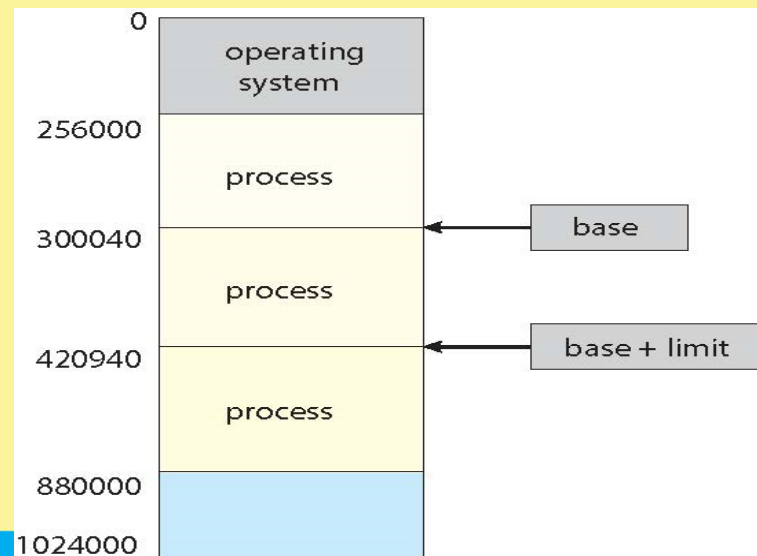
## Background (Cont.)

- Completing a memory access may take many cycles of the CPU clock. In such a case the processor needs to **stall** since it does not have the data required to complete the instruction it is executing.
- **Cache** sits between main memory and CPU registers to deal with the “stall” issue.
- Protection of memory is required to ensure correct operation:
  - User process cannot access OS memory
  - One user process cannot access the memory of another user process.

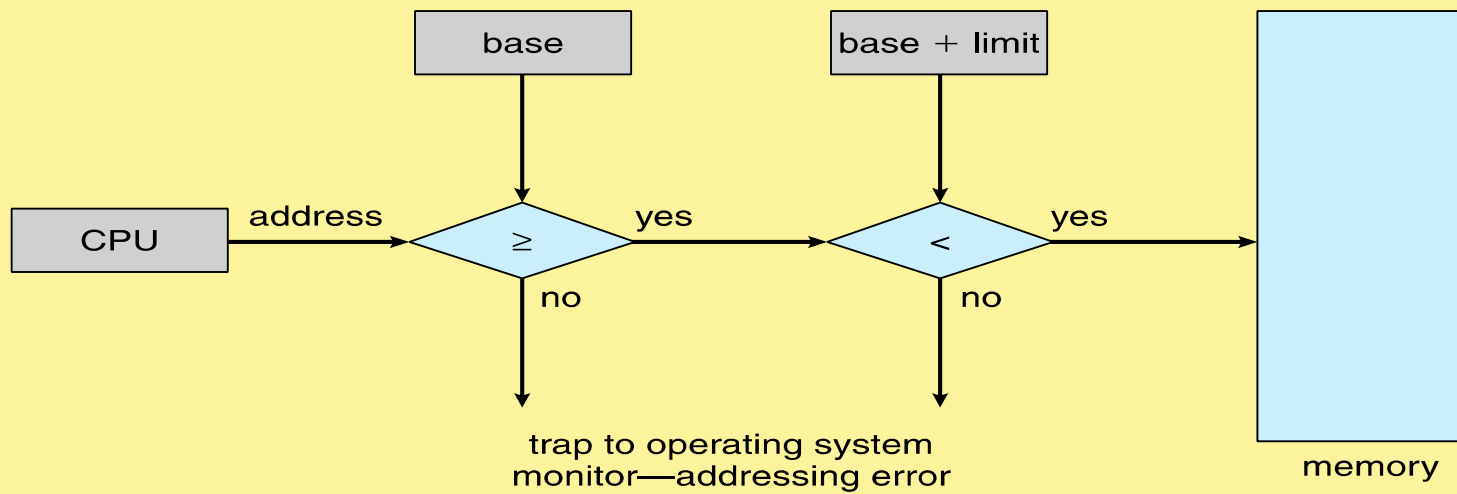


# Memory Protection

- A **base register** (holding the smallest legal physical address of a program in memory) and a **limit register** (specifies the size of the program) define the boundary of a program in memory.
- CPU must check that every memory access generated in user mode is between the base and base+limit for that user



# Hardware Address Protection





# Address Binding

- A program residing on the disk needs to be brought into memory in order to execute. Such a program is usually stored as a binary executable file and is kept in an **input queue**.
- In general, we do not know a priori where the program is going to reside in memory. Therefore, it is convenient to assume that the first physical address of a program always starts at location 0000.
- Without some hardware or software support, program must be loaded into address 0000
- It is impractical to have first physical address of user process to always start at location 0000.
- Most (all) computer systems provide hardware and/or software support for memory management



# Address Binding (Cont.)

- In general, addresses are represented in different ways at different stages of a program's life
  - Addresses in the source program are generally symbolic
    - For example, variable "count"
  - A compiler typically **binds** these symbolic addresses to relocatable addresses
    - For example, "14 bytes from beginning of this module"
  - Linker or loader will bind relocatable addresses to absolute (physical) addresses
    - For example, 74014
  - Each binding maps one address space to another address space



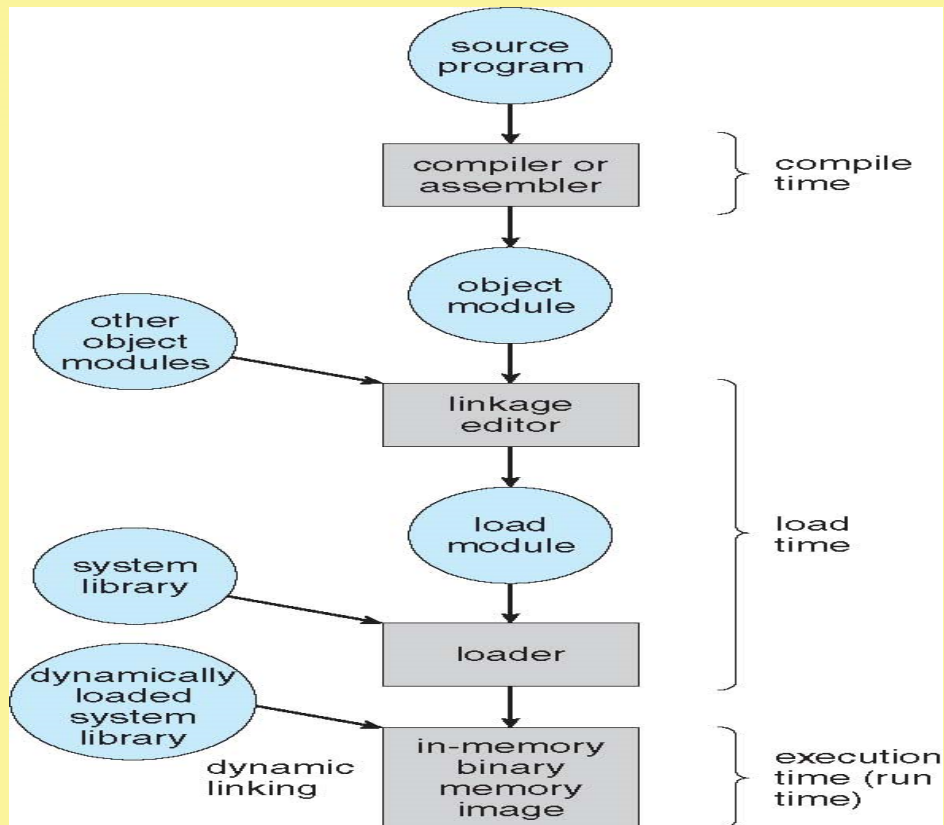
## Binding of Instructions and Data to Memory

Address binding of instructions and data to memory addresses can happen at three different points in time:

- **Compile time:** If memory location known a priori, **absolute code** can be generated; must recompile code if starting location changes.
- **Load time:** If memory location is not known at compile time and no hardware support is available, **relocatable code** must be generated (in software).
- **Execution time:** Binding delayed until run time if the process can be moved during its execution from one memory segment to another
  - Need hardware support for address maps (e.g., base and limit registers)



## Multistep Processing of a User Program



# Logical vs. Physical Address Space


- The concept of a **logical address space** that is bound to a separate **physical address space** is central to proper memory management
  - **Logical address** – generated by the CPU.
  - **Physical address** – address seen by the memory unit
- Logical and physical addresses are:
  - The same in compile-time and load-time address-binding schemes;
  - They differ in execution-time address-binding scheme. In that case the logical address is referred to as **virtual address**.

We use Logical address and virtual address interchangeably
- **Logical address space** is the set of all logical addresses generated by a program
- **Physical address space** is the set of all physical addresses corresponding to a given logical address space.



# Memory-Management Unit (MMU)


- Hardware device that at run time maps virtual addresses to physical address



- Many methods possible, covered in the rest of this chapter
- The user program deals with *logical* addresses; it never sees the *real* physical addresses
  - Execution-time binding occurs when reference is made to location in memory
  - Logical address bound to physical addresses

## Dynamic relocation using a relocation register

- To start, consider simple scheme where the value in the base register is added to every address generated by a user process at the time it is sent to memory
  - Base register now called **relocation register**
  - MS-DOS on Intel 80x86 used 4 relocation registers



# Dynamic Loading

- Until now we assumed that the entire program and data has to be in main memory to execute
- **Dynamic loading** allows a routine (module) to be loaded into memory only when it is called (used)
- Results in better memory-space utilization; an unused routine is never loaded
- All routines kept on disk in relocatable load format
- Useful when large amounts of code are needed to handle infrequently occurring cases (e.g., exception handling)
- No special support from the operating system is required
  - It is the responsibility of the users to design their programs to take advantage of such a method
  - OS can help by providing libraries to implement dynamic loading





# Dynamic Linking

- **Dynamically linked libraries** – system libraries that are linked to user programs when the programs are run.
  - Similar to dynamic loading. But, linking rather than loading is postponed until execution time
- Small piece of code, **stub**, used to locate the appropriate memory-resident library routine
- Stub replaces itself with the address of the routine, and executes the routine
- Operating system checks if routine is in processes' memory address
  - If not in address space, add to address space
- Dynamic linking is particularly useful for libraries
- System also known as **shared libraries**



# Contiguous Allocation

- Main memory must support both OS and user processes
- Limited resource -- must allocate efficiently
- Contiguous allocation is one early method
- Main memory is usually divided into two **partitions**:
  - Resident operating system, usually held in low memory with interrupt vector
  - User processes are held in high memory
  - Each process contained in single contiguous section of memory

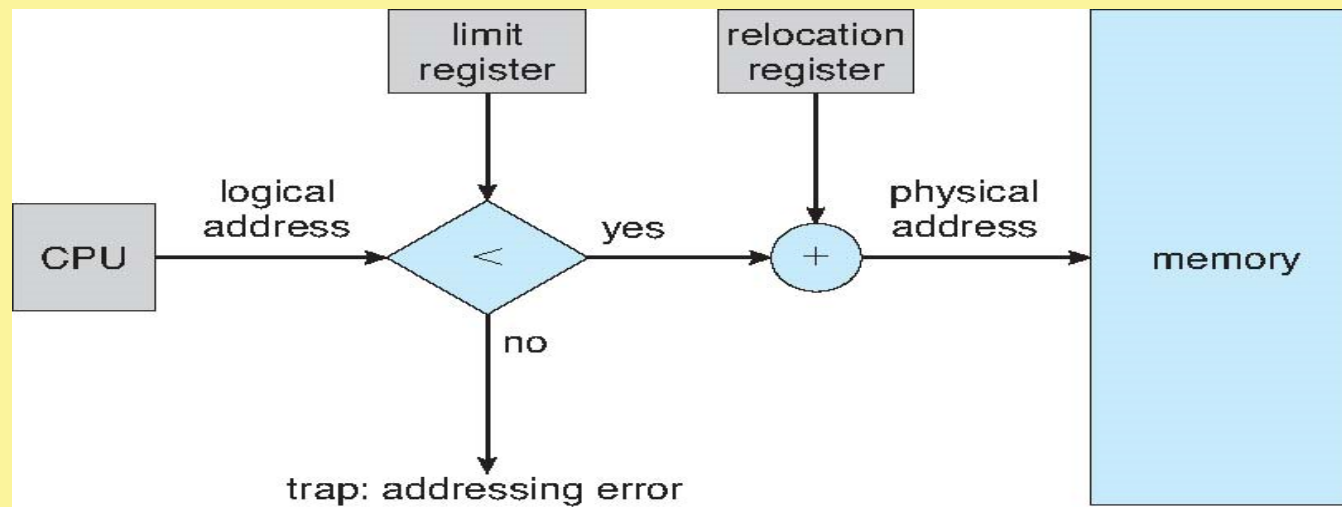


## Contiguous Allocation (Cont.)

- Relocation registers used to protect user processes from each other, and from changing operating-system code and data
  - Base register contains value of smallest physical address
  - Limit register contains range of logical addresses – each logical address must be less than the limit register
  - MMU maps logical address *dynamically*
  - Can then allow actions such as kernel code being **transient** – comes and goes as needed. Thus, kernel can change size dynamically.

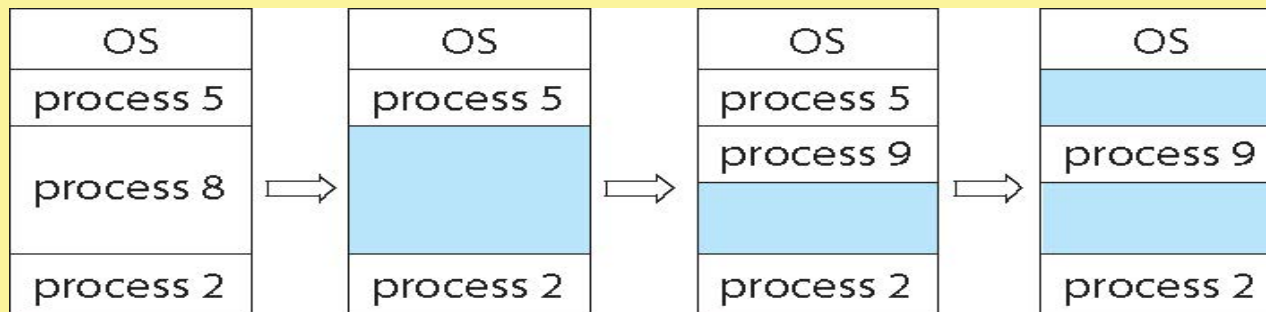


## Hardware Support for Relocation and Limit Registers



# Multiple-partition allocation

- **Variable-partition** -- sized to a given process' needs.
- **Hole** – block of available memory; holes of various size are scattered throughout memory
- When a process arrives, it is allocated memory from a hole large enough to accommodate it
- Process exiting frees its partition, adjacent free partitions combined
- Operating system maintains information about:  
a) allocated partitions   b) free partitions (holes)



# Dynamic Storage-Allocation Problem

- How to satisfy a request of size  $n$  from a list of free holes?
  - **First-fit**: Allocate the **first** hole that is big enough
  - **Best-fit**: Allocate the **smallest** hole that is big enough; must search entire list, unless the list is ordered by size.
    - Produces the smallest leftover hole
  - **Worst-fit**: Allocate the **largest** hole; must also search entire list, unless the list is ordered by size
    - Produces the largest leftover hole
- First-fit and best-fit are better than worst-fit in terms of speed and storage utilization



# Fragmentation

- **External Fragmentation** – total memory space exists to satisfy a request, but it is not contiguous and therefore cannot be used.
  - First fit analysis reveals that given  $N$  allocated blocks, another  $0.5 N$  blocks will be lost to fragmentation
    - 1/3 of memory may be unusable -> **50-percent rule**
- **Internal Fragmentation** – allocated memory may be slightly larger than requested memory.
  - Can happen if there is hole of size 15,000 bytes and a process needs 14,900 bytes; Keeping a hole of size 100 bytes is not worth the effort so the process is allocated 15,000 bytes.
  - The size difference of 100 bytes is memory internal to a partition, but not being used



# Fragmentation (Cont.)

Reduce external fragmentation by **compaction**

- Shuffle memory contents to place all free memory together in one large block
- Compaction is possible *only* if relocation is dynamic, and is done at execution time
- I/O problem -- cannot perform compaction while I/O is in progress involving memory that is being compacted.
  - Latch job in memory while it is involved in I/O
  - Do I/O only into OS buffers





# Non-contiguous Allocation

- Partition the a program into a number of small units, each of which can reside in a different part of the memory.
- Need hardware support.
- Various methods to do the partitions:
  - Segmentation.
  - Paging
  - paged segmentation.




# Segmentation


- Memory-management scheme that supports user's view of memory
- A program is a collection of segments -- a logical unit such as:  
main program, procedure, function, method,  
object, local variables, global variables,  
common block, stack, symbol table, arrays
- Each segment can do reside in different parts of memory.
- Way to circumvent the contiguous allocation requirement.



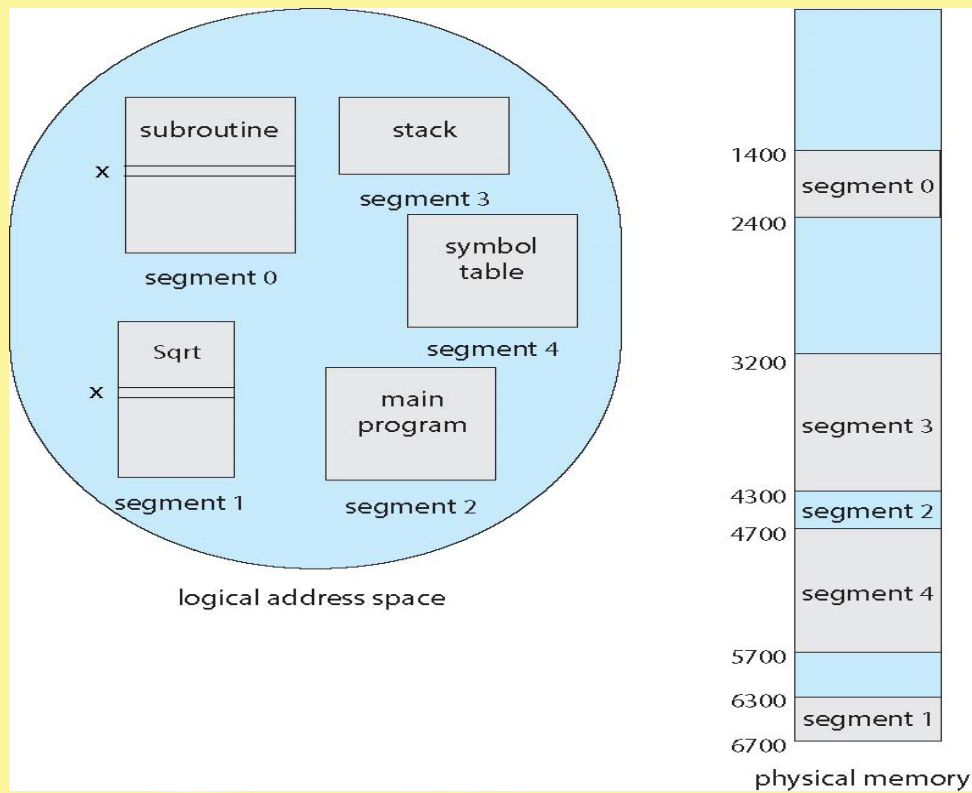
# User's View of a Program



# Two Dimensional Addresses



# Logical and Physical Memory

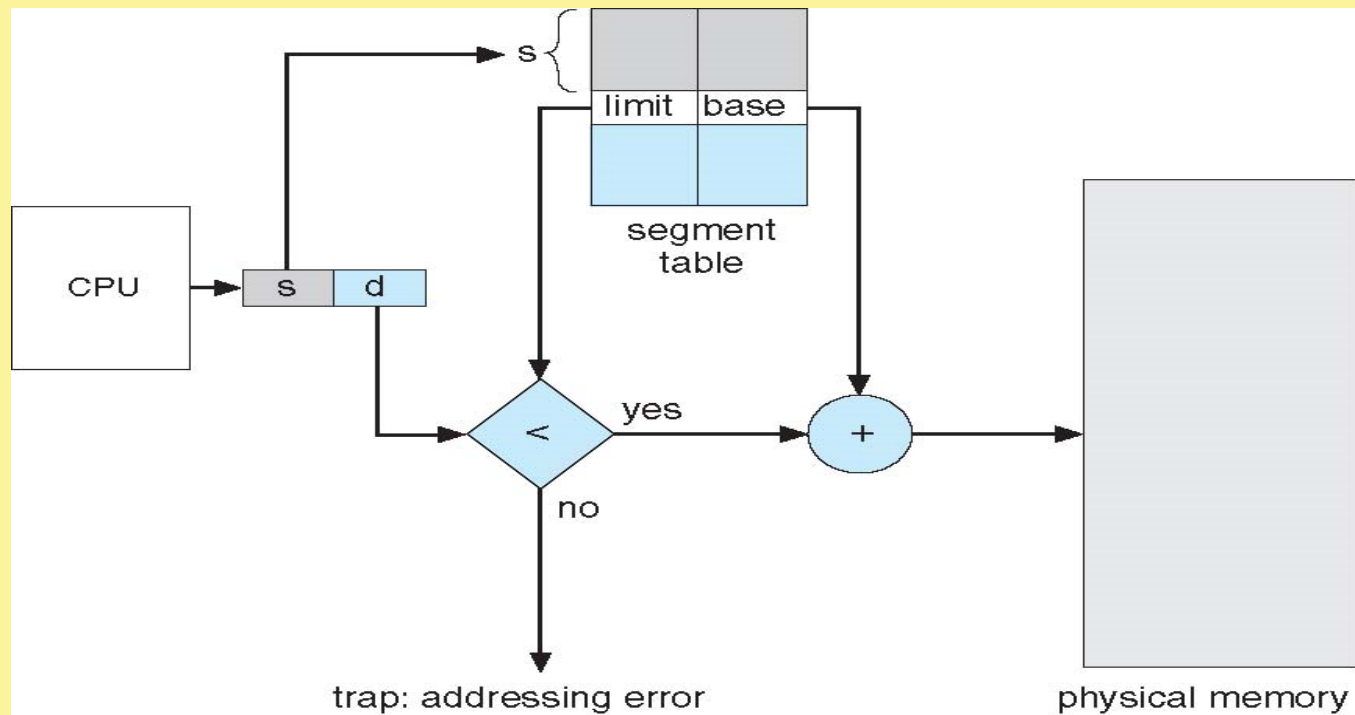


# Segmentation Architecture

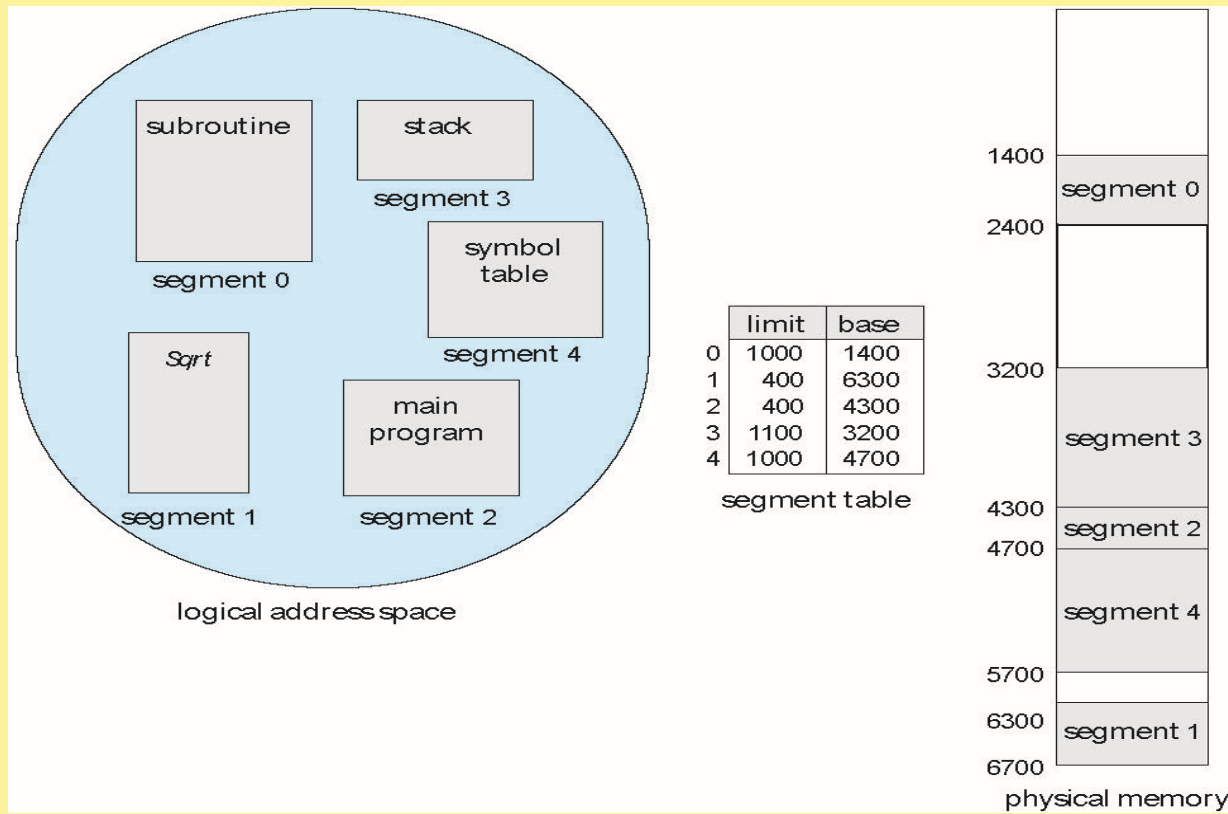
- Logical address consists of a two tuple:  
    <segment-number, offset>
- Need to map a two-dimensional logical addresses to a one-dimensional physical address. Done via **Segment table**:
  - **base** – contains the starting physical address where a segments reside in memory
  - **limit** – specifies the length of the segment
- Segment table is kept in memory
  - **Segment-table base register (STBR)** points to the segment table's location in memory
  - **Segment-table length register (STLR)** indicates number of segments used by a program;  
    segment number **s** is legal if **s** < **STLR**



# Segmentation Hardware



# Example of Segmentation





# Paging

- Physical address space of a process can be non-contiguous.
- Process is divided into fixed-size blocks, each of which may reside in a different part of physical memory.
- Divide physical memory into fixed-sized blocks called **frames**
  - Size of a frame is power of 2 between 512 bytes and 16 Mbytes
- Divide logical memory into blocks of same size as frames called **pages**
- Backing store (dedicated disk), where the program is permanently residing, is also split into storage units (called **blocks**), which are the same size as the frame and pages.
- Physical memory allocated whenever the latter is available
  - Avoids external fragmentation
  - Still have Internal fragmentation



## Paging (Cont.)

- Keep track of all free frames
- To run a program of size ***N*** pages, need to find ***N*** free frames and load program from backing store.
- Set up a **page table** to translate logical to physical addresses
- Page table is kept in memory.
  - **Page-table base register (PTBR)** points to the page table
  - **Page-table length register (PTLR)** indicates size of the page table
- Still have Internal fragmentation



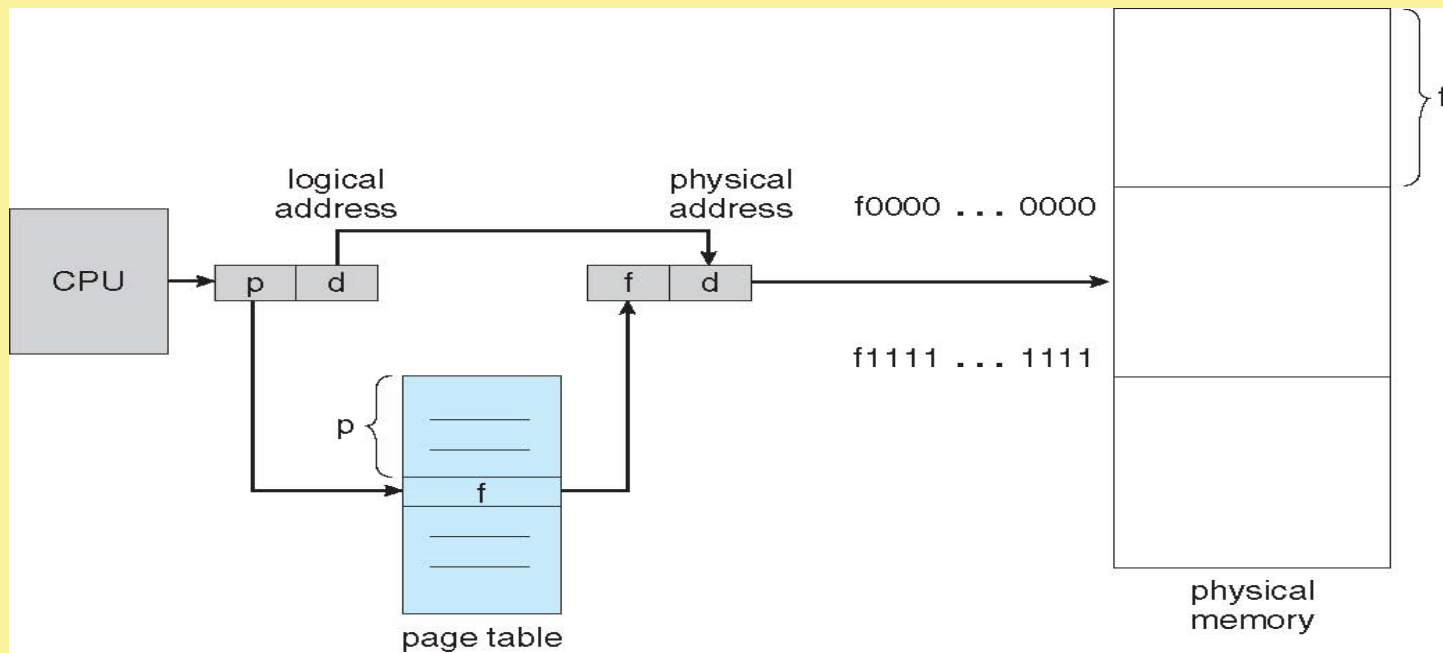
# Address Translation Scheme

- Assume the logical address space is  $2^m$ . (How is  $m$  determined?)
- Assume page size is  $2^n$
- Address generated by CPU is divided into:
  - **Page number** ( $p$ ) – used as an index into a **page table** which contains base address of each page in physical memory. Size of  $p$  is “ $m - n$ ”
  - **Page offset** ( $d$ ) – combined with base address to define the physical memory address that is sent to the memory unit. Size of  $d$  is “ $n$ ”.

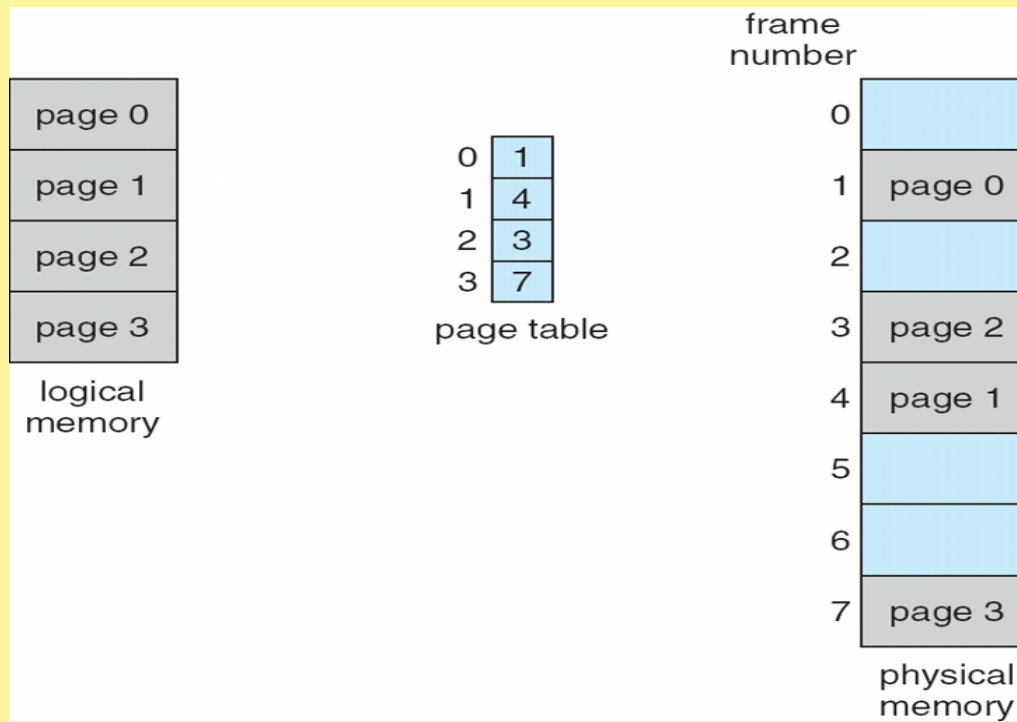
| page number | page offset |
|-------------|-------------|
| $p$         | $d$         |
| $m - n$     | $n$         |



# Paging Hardware

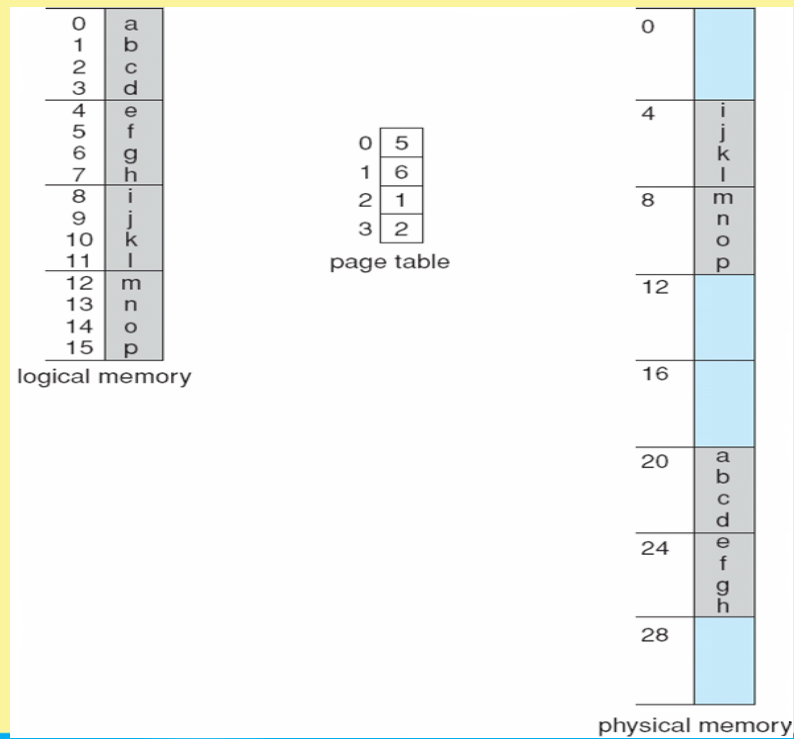


## Paging Model of Logical and Physical Memory



# Paging Example

Assume  $m = 4$  and  $n = 2$  and 32-byte memory and 4-byte pages



# Internal Fragmentation

- Calculating internal fragmentation
  - Page size = 2,048 bytes
  - Process size = 72,766 bytes
  - 35 pages + 1,086 bytes
  - Internal fragmentation of  $2,048 - 1,086 = 962$  bytes
  - Worst case fragmentation = 1 frame – 1 byte
  - On average fragmentation = 1 / 2 frame size
  - So small frame sizes desirable?
  - But each page table entry takes memory to track
  - Page sizes growing over time
    - Solaris supports two page sizes – 8 KB and 4 MB
- By implementation process can only access its own memory

